

EGN4060c Intro to Robotics

Lecture 9: Machine Learning: Supervised Classifiers

Instructor: Dr. Gita Sukthankar
Email: gitaras@eecs.ucf.edu

Upcoming Due Dates

- Much happening in the next few weeks!
- Extended date for lab 2 report due on Wed; no further extensions
- Homework 2 (architectures) due on Wed
- Lab 3 report due next Wed Oct 1st
- Midterm exam: Wed Oct 8th (exam review on the 6th)

Wavefront Summary

- Maintain 2 data structures (also can track current cost)
 - Old cost map
 - New cost map
- Iterate over the grid until the start square has a non-zero value
- For each cell:
 - Find lowest cost neighboring, unoccupied square and add 1 to the cost
 - If the current cost is 0 or if the new cost is lower than the current cost, annotate the cell (new cost map) with the new cost.
- Your path is defined by any uninterrupted sequence of decreasing numbers reaching the goal

Wavefront (Complete)

- You're done
 - Remember, 0's should only remain if unreachable regions exist

| | | | | | | | | | | | | | | | |
|---|----|----|----|----|----|----|----|----|----|---|----|----|----|----|----|
| 7 | 18 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 9 | 9 | 9 | 9 | 9 |
| 6 | 17 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 8 | 8 | 8 | 8 |
| 5 | 17 | 16 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 7 | 7 | 7 |
| 4 | 17 | 16 | 15 | 15 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 6 | 6 | 6 |
| 3 | 17 | 16 | 15 | 14 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 5 | 5 | 5 |
| 2 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 4 |
| 1 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 |
| 0 | 17 | 16 | 15 | 14 | 13 | 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 2 |
| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |

Following the Path

- To find the shortest path, according to your metric, simply always move toward a cell with a lower number
 - The numbers generated by the Wavefront planner are roughly proportional to their distance from the goal

Two
possible
shortest
paths
shown



Java: Bounds Checking

- Each array object has a public constant called `length` that stores the size of the array
- It is referenced using the array name:

`scores.length`

- Note that `length` holds the number of elements, not the largest index

Java: Initializer Lists

- An *initializer list* can be used to instantiate and fill an array in one step
- The values are delimited by braces and separated by commas
- Examples:

```
int[] units = {147, 323, 89, 933, 540,  
               269, 97, 114, 298, 476};
```

```
char[] letterGrades = {'A', 'B', 'C', 'D', 'F'};
```

Two-Dimensional Arrays

- A two-dimensional array is declared by specifying the size of each dimension separately:

```
int[][] scores = new int[12][50];
```

- A array element is referenced using two index values:

```
value = scores[3][6]
```

- The array stored in one row can be specified using one index

MapGUI: class for map display

```
MapGUI map = new MapGUI();  
map.getMap();  
map.moveRobot(row, column, angle);  
map.moveRobot(row, column, direction);  
map.getRobotLocation();  
Example map file format: 0=empty, 1=wall, 2=goal
```

```
6 8  
00100010  
00000000  
00000100  
01101100  
00011000  
10010002  
2 1 120
```

Overview

- You've learned how to write scripts for the robot.
- You've learned how to have the robot create and execute a plan based on a pre-specified world model.
- But what if you want your robot to be able to learn from experiences?
- The answer----machine learning!

Problem Formulation

- Input: robot sensor reading
- Output: instructions for robot
- Additional information:
 - Example pairs of input and output (supervised learning)
 - Fitness function (evolutionary computing)
 - Reward function (reinforcement learning)
- For the next lab, you' ll be doing reinforcement learning.

Classification

- Assign object/event to one of a given finite set of categories.
 - Medical diagnosis
 - Credit card applications or transactions
 - Fraud detection in e-commerce
 - Worm detection in network packets
 - Spam filtering in email
 - Recommended articles in a newspaper
 - Recommended books, movies, music, or jokes
 - Financial investments
 - DNA sequences
 - Spoken words
 - Handwritten letters
 - Astronomical images

Computer vision and machine learning are closely related disciplines.

Planning / Control

- Performing actions in an environment in order to achieve a goal.
 - Solving calculus problems
 - Playing checkers, chess, or backgammon
 - Balancing a pole
 - Driving a car or a jeep
 - Flying a plane, helicopter, or rocket
 - Controlling an elevator
 - Controlling a character in a video game
 - Controlling a mobile robot

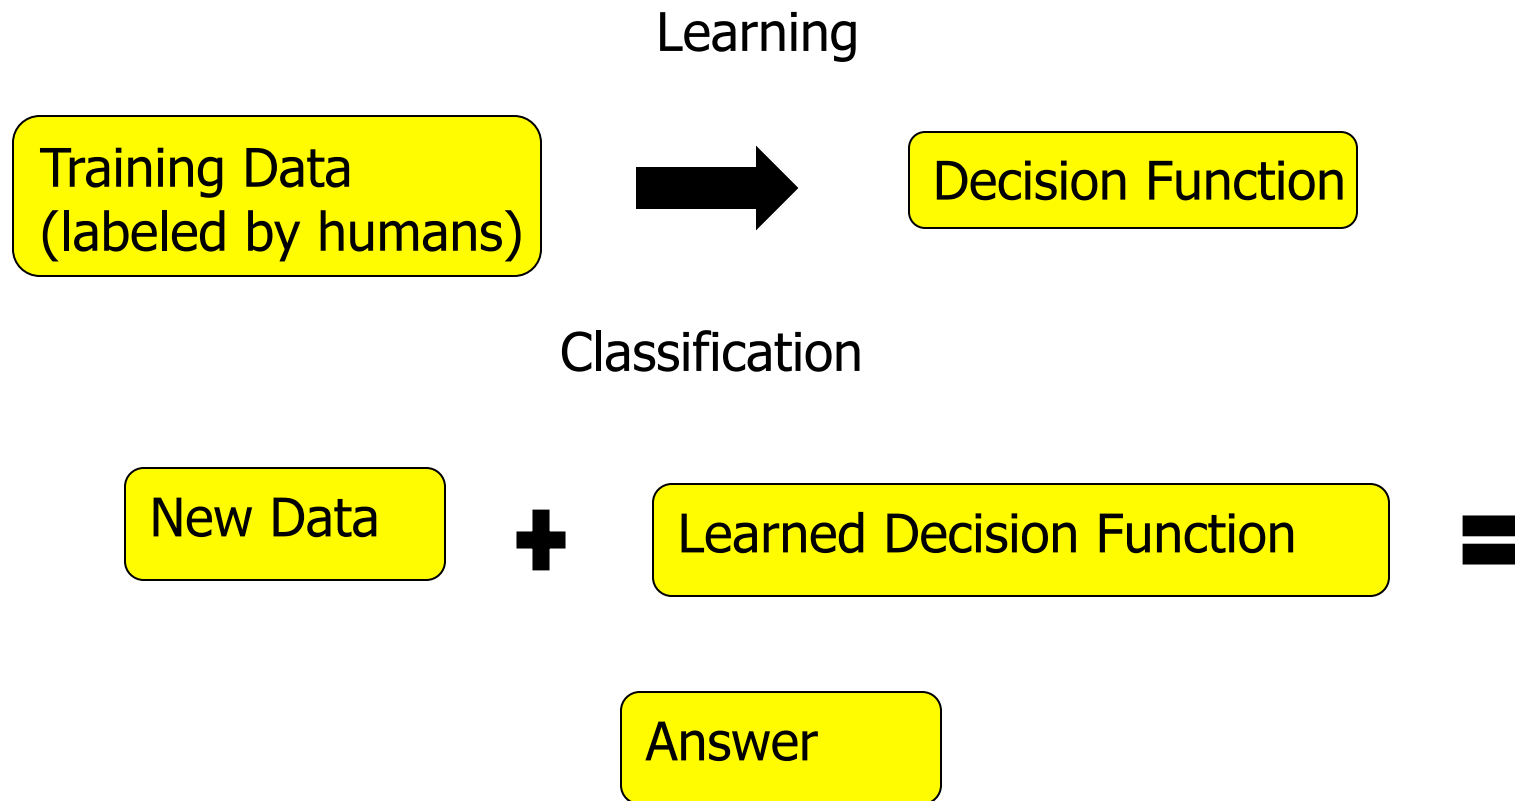
Measuring Performance

- Classification accuracy
- Solution correctness
- Solution quality (length, efficiency)
- Speed of performance

The Time is Ripe for ML!

- Many basic effective and efficient algorithms available.
- Large amounts of on-line data available.
- Large amounts of computational resources available.

Supervised Classifiers



How to learn the decision function?

Regression

- Problem: from set of exemplars and known x , predict value for unknown y
- Assume a model for the data based on parameters $y=f(x;p)$
 - x is data
 - p are parameters
 - $f(x)$ is the model
- For example:
 - You could model your data with a linear model:
 $f(x;p)=ax+b$
 - The parameters of this model are: $p=(a,b)$

Regression Basics

- Learning: estimate the model parameters p from training data (x, y)
- For $f(x; p) = a$

$$a = \frac{1}{N} \sum_i y_i$$

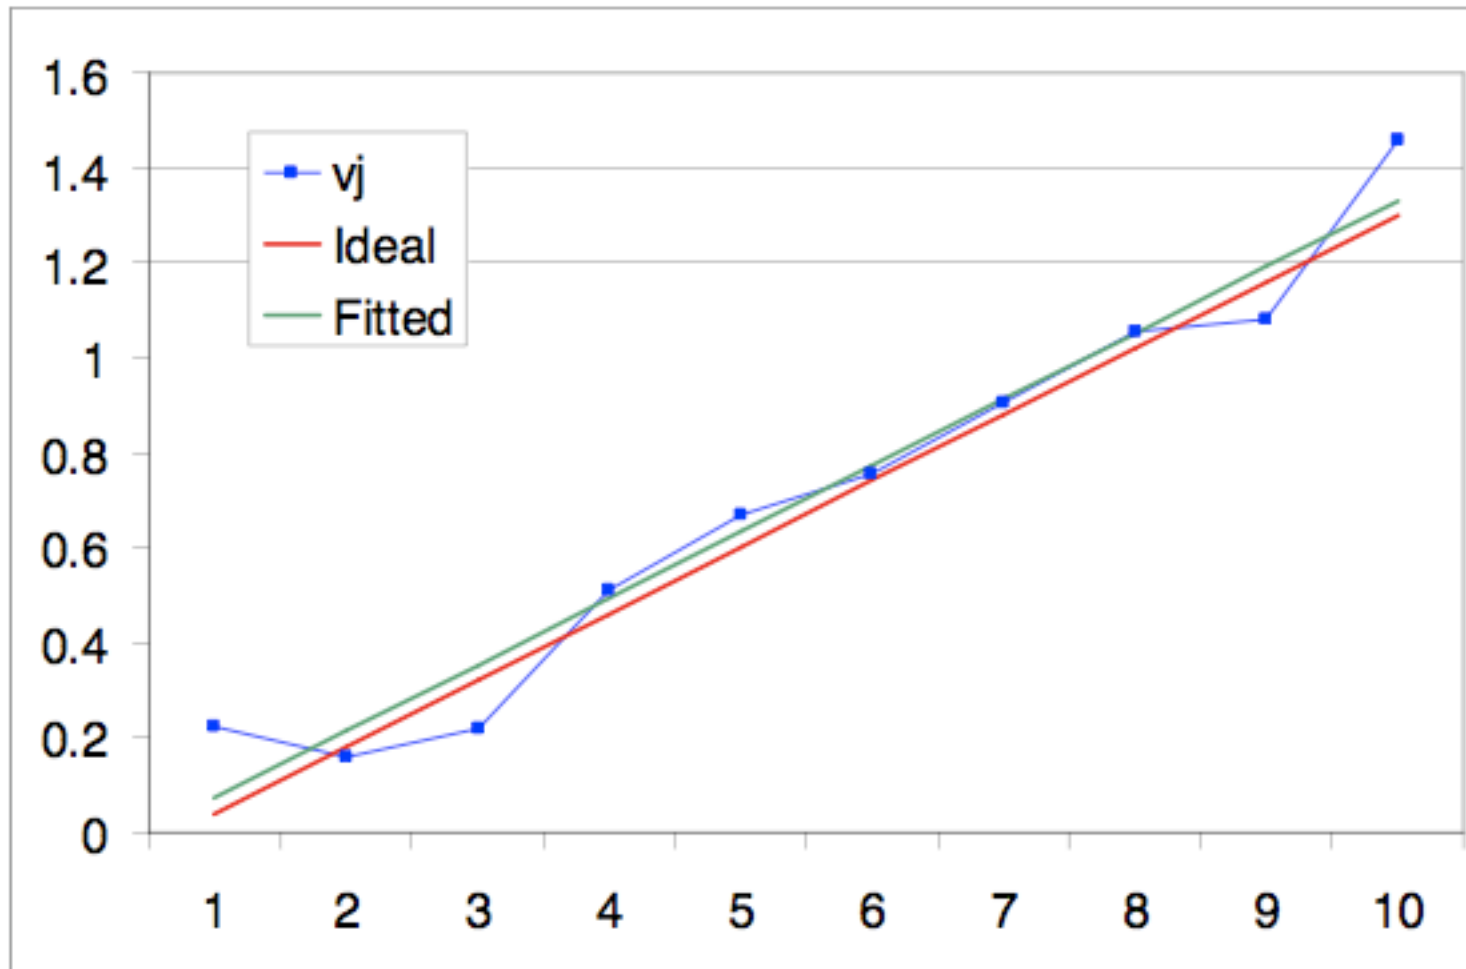
Linear Regression

- Linear model: $f(x;p)=ax+b$
- Ordinary least squares method: minimize the residual which can be done in a closed form way

$$\begin{aligned} S_x &= n \cdot \bar{x} = \sum_{j=1}^n x_j & S_y &= n \cdot \bar{y} = \sum_{j=1}^n y_j \\ S_{xx} &= \sum_{j=1}^n x_j^2 & S_{xy} &= \sum_{j=1}^n x_j y_j \end{aligned}$$

$$a = \frac{nS_{xy} - S_x S_y}{nS_{xx} - (S_x)^2} \quad b = \frac{1}{n} (S_y - aS_x)$$

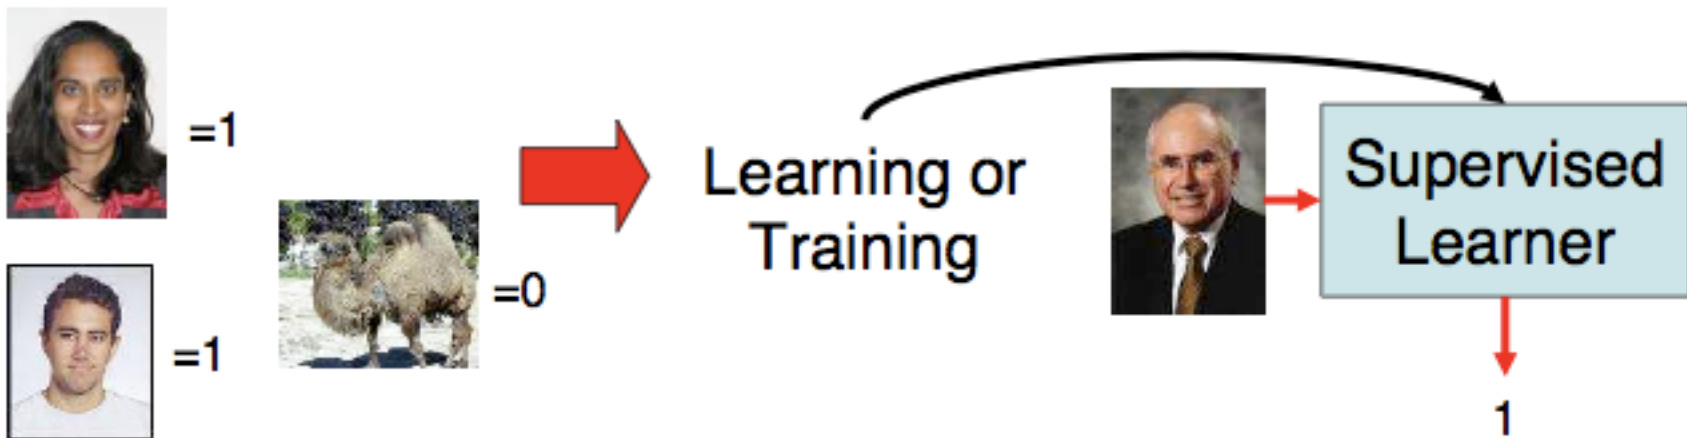
Graph the Result



Once the parameters have been “learned” they can be used to predict the answers for unseen values.

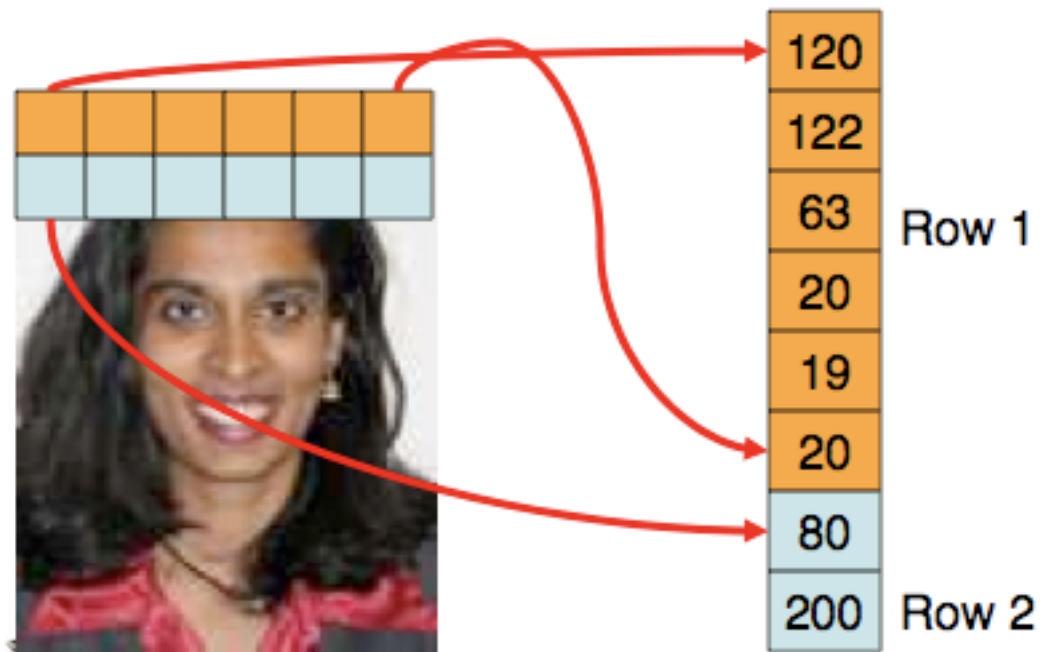
Supervised Learning

- Given training data sequence of inputs x and outputs y $(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)$
- Learn to predict the output y
- Input x can be real, discrete, or multi-dimensional



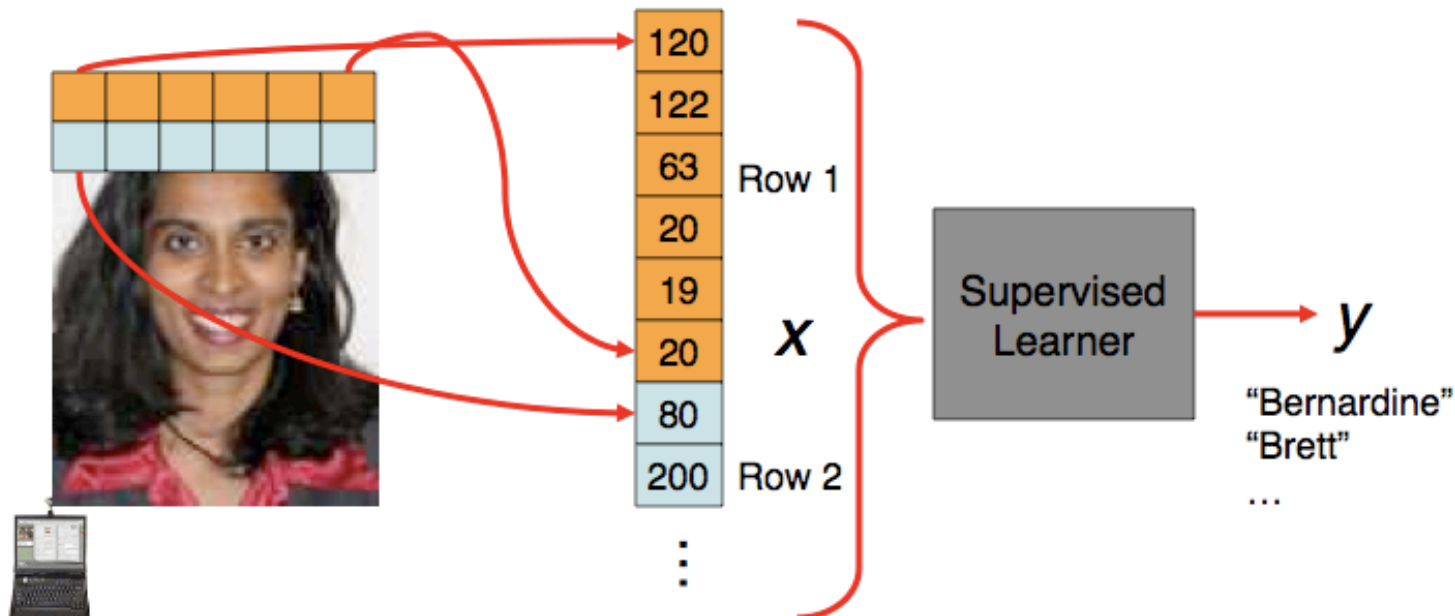
Face Detection

- X is a vector of pixel intensities



Supervised Learning

Y is a discrete output class just binary



Supervised Learning

- Labeled examples
- Real input, with desired output



“Bernardine”



“Unknown”



“Brett”

Binary Classification

- Output has 2 classes: true or not true
 - Usually defined as $\{0,1\}$ or $\{-1,+1\}$



“Face” = 1



“Not face” = -1



“Face” = 1

Multi-Class Classification

- Output has many classes
- Usually defined as $\{1....k\}$



“Bernardine” = 1



“Camel” = 3



“Brett” = 2

Regression and Classification

- Regression and classification are very similar
- We can often define a binary classification as

$$y = \text{sgn}(f(x))$$
$$y = \begin{cases} 1 & f(x) > 0 \\ -1 & f(x) \leq 0 \end{cases}$$

Classification Approaches

- Nearest neighbor
- Naïve Bayes
- Decision trees
- Neural networks
 - Perceptrons
 - Multi-layer perceptron
- All of these different methods answer the same question.
 - Given a known input and training data, what is the unknown class label of a new example?

Definitions

- C is the set of classes, for binary classification
 - $C=\{-1,+1\}$
- X is the input space
 - Typically a real d-dimensional vector

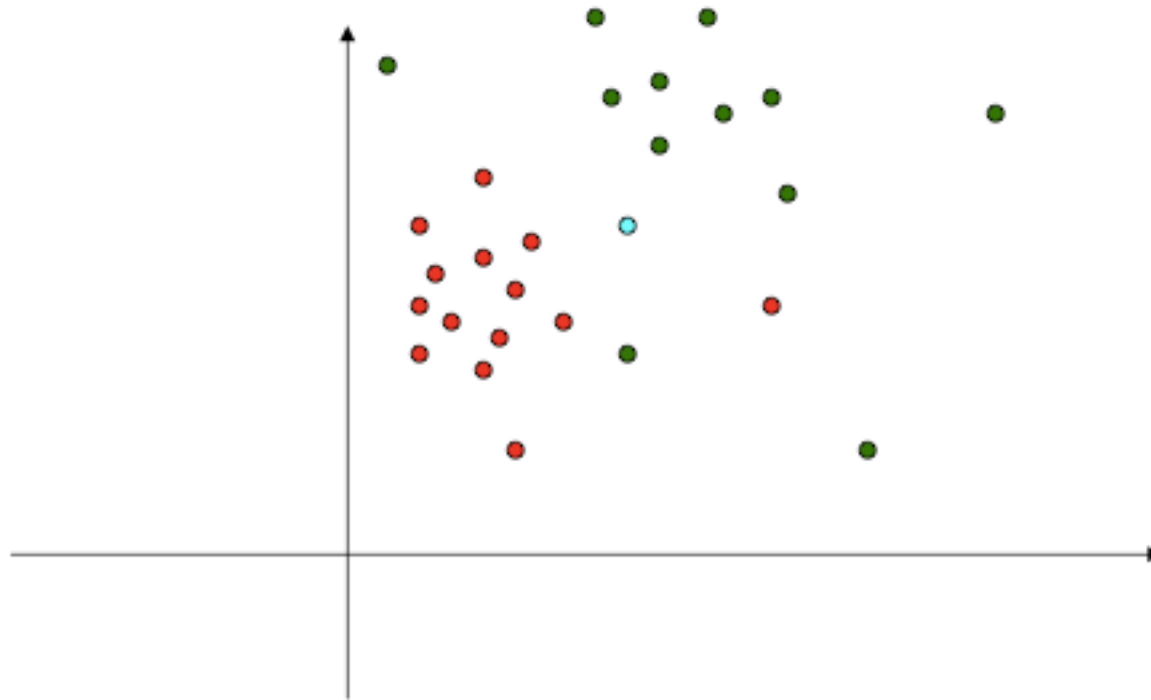
$$x \in \mathbb{R}^D$$

- Y is the output so

$$y \in C$$

Problem

- Green dots are one class, red dots are the other
- Given an example (blue) which class does it belong to?



Nearest Neighbor

- Simple idea:
 - Look for closest example x in training data and use its output y as the output
- Mathematically

$$i^* = \arg \min_i (D(x, x_i))$$

$$y = y_{i^*}$$

Similar in concept to k-means clustering (which is an unsupervised learning technique)

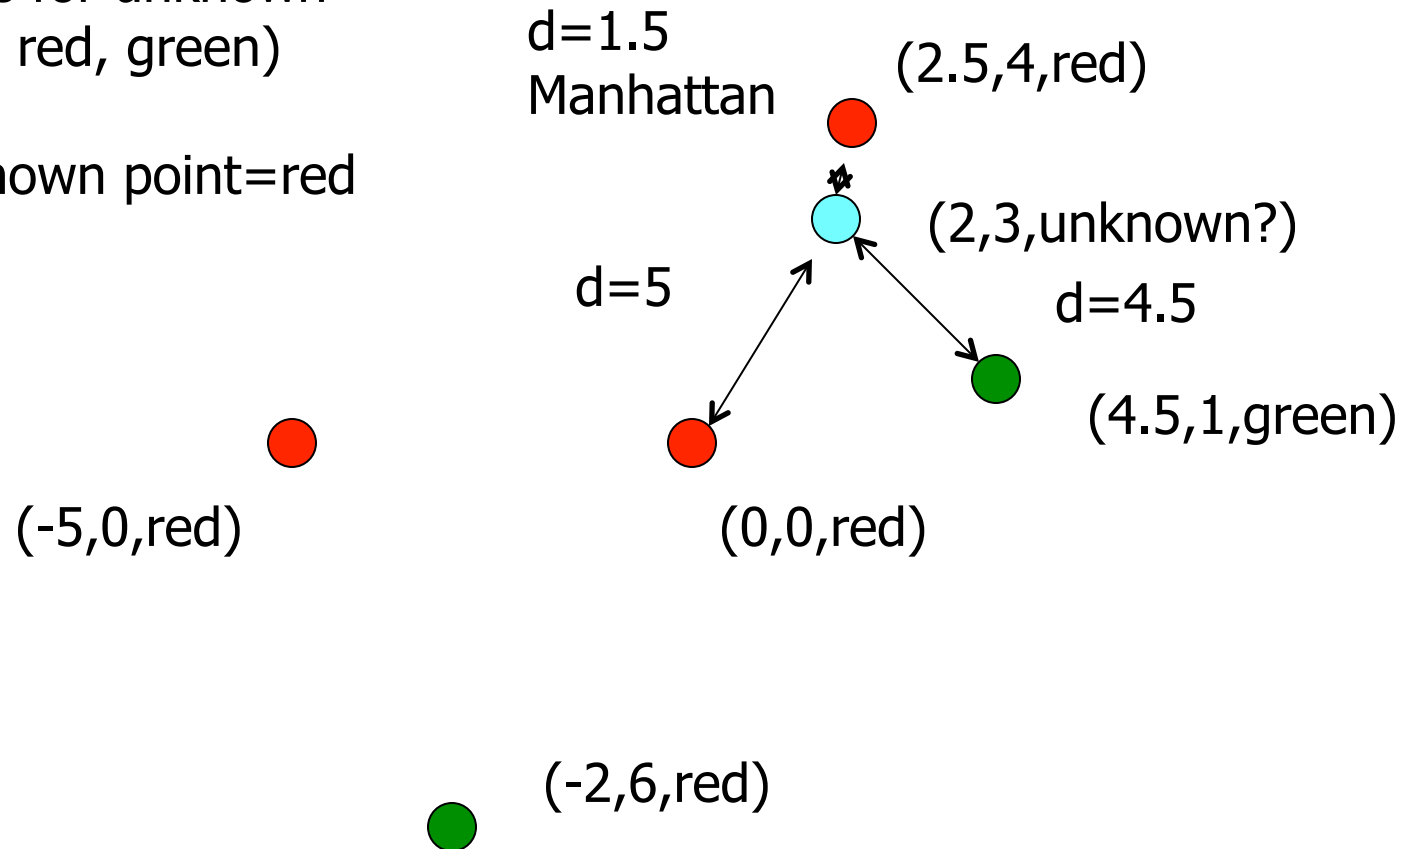
Distance Function

- Euclidean $D(a,b) = \sqrt{\sum_j (a_j - b_j)^2}$
- Manhattan $D(a,b) = \sum_j |a_j - b_j|$
- Choice in distance function affects answer

3 Nearest Neighbor

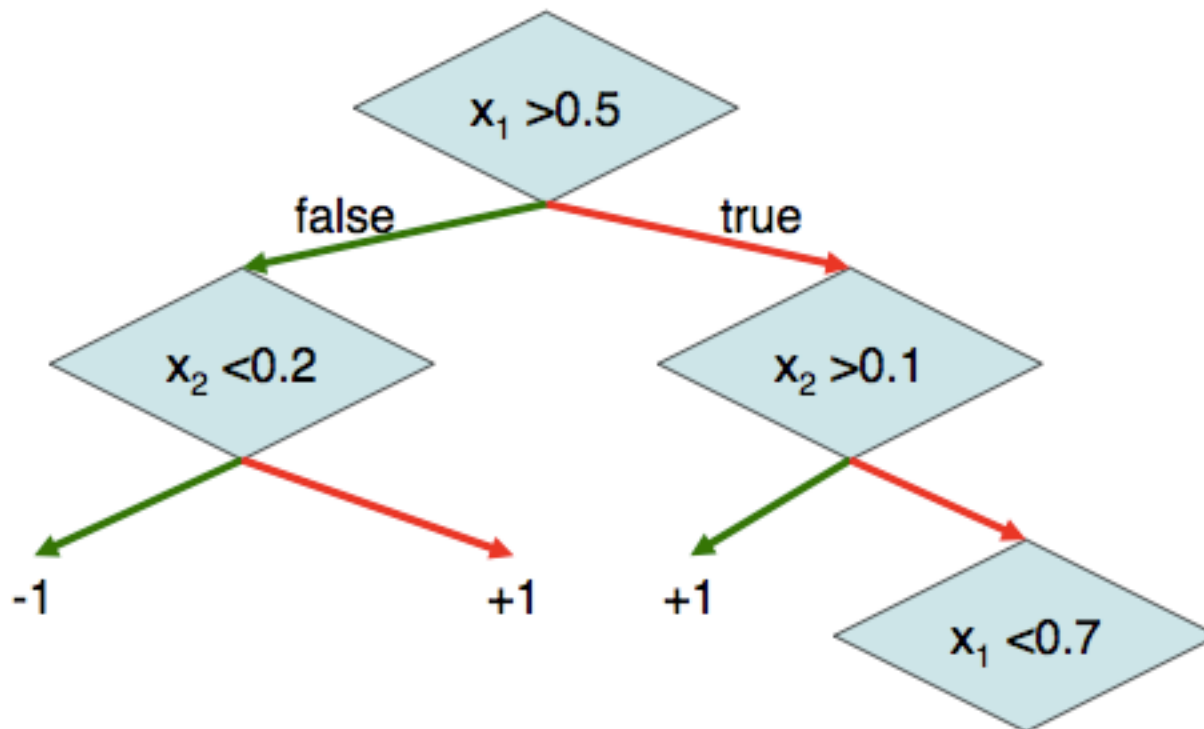
Votes for unknown
(red, red, green)

Unknown point=red



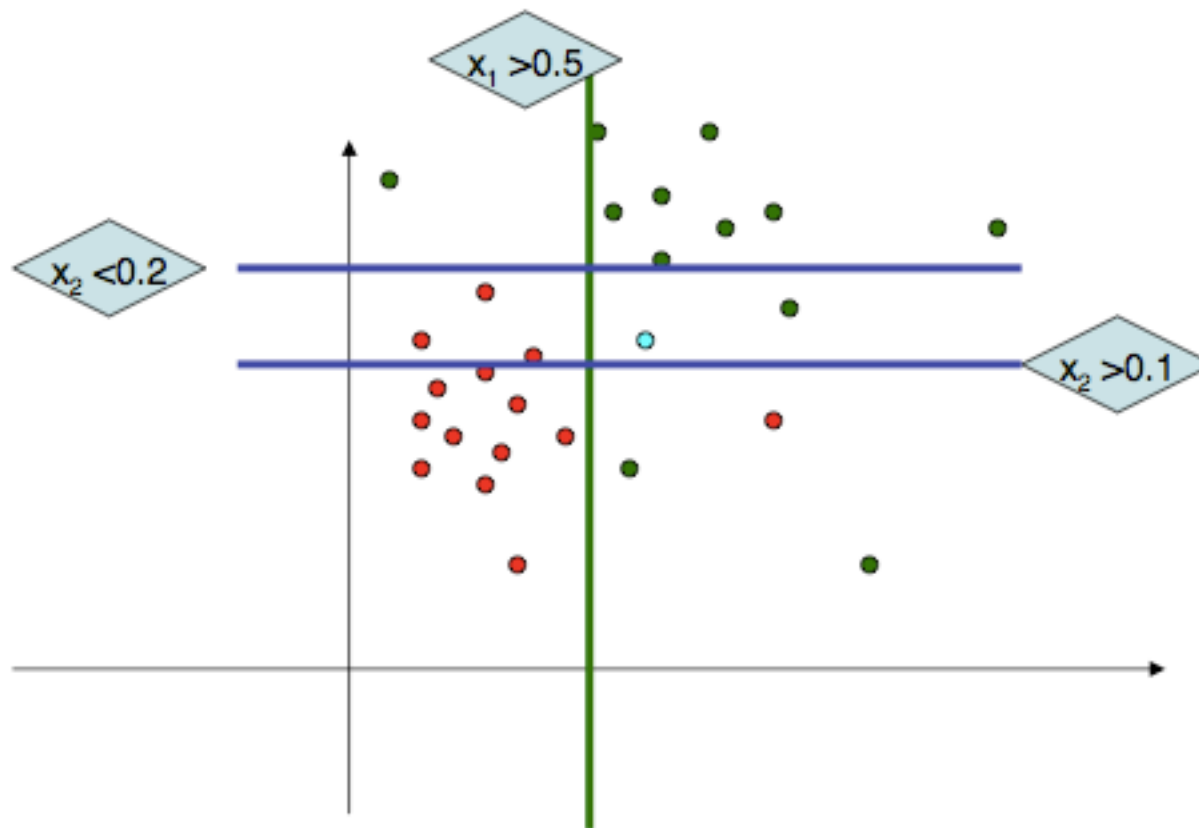
Decision Trees

- Basic concept:
 - Learn a nested set of if/else rules
 - Tree of decisions or splitting points



Decision-Tree

- Axis aligned splitting planes



Decision Trees

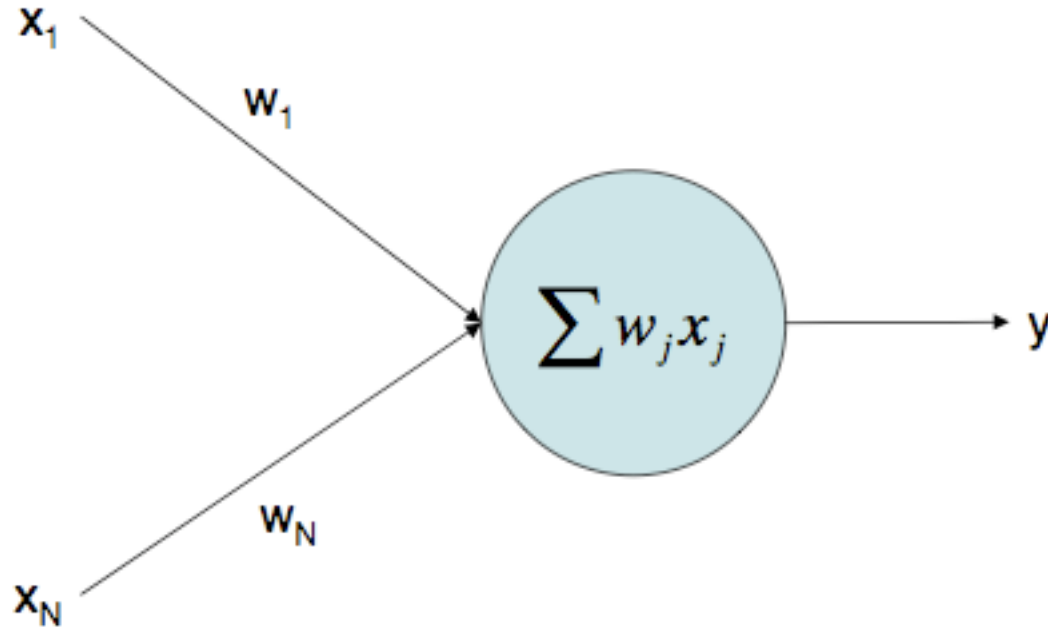
- Best example is C4.5 algorithm
 - Works on discrete and continuous data sets
- Key idea to building trees
 - Deciding where to introduce a split and when to start splitting

Neural Networks

- Modeled on biological networks
- Activation related to strength of outputs
- First proposed in 1940s by McCulloch and Pitts
- Shot down by Minsky around 1969
- Returned to popularity in 1980s with the backpropagation algorithm

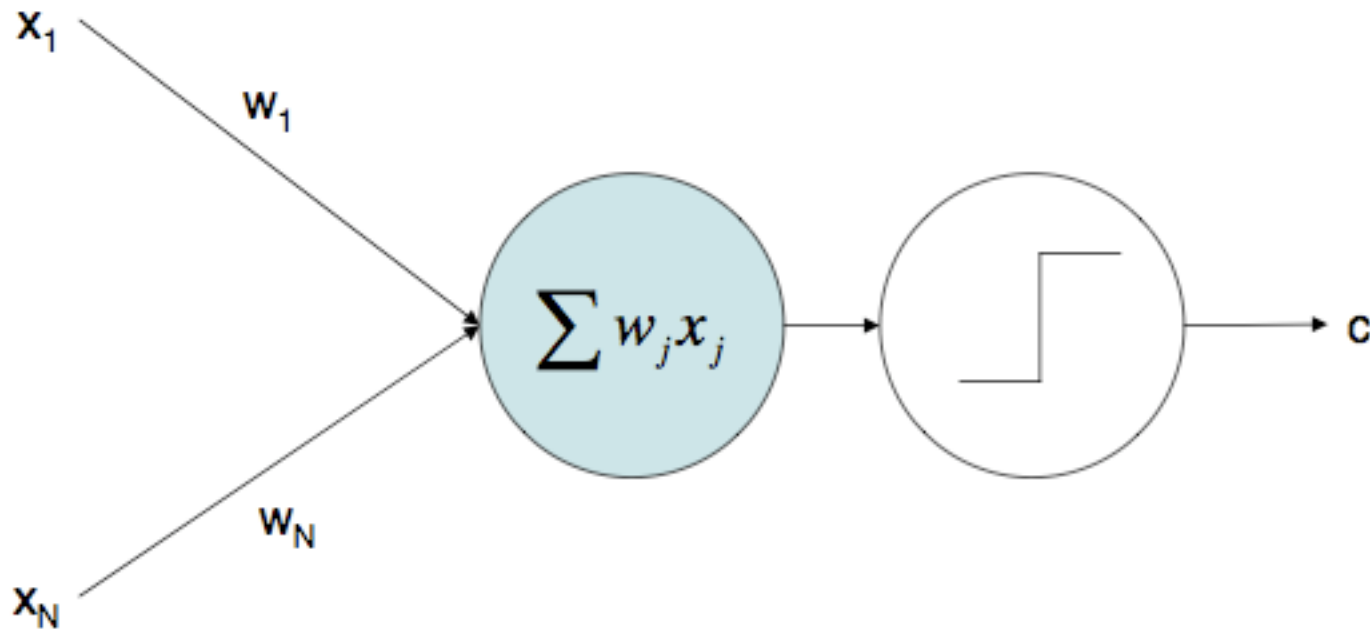
Perceptron Learning

- Single linear “neuron”



Perceptrons for Classification

- Output through step function



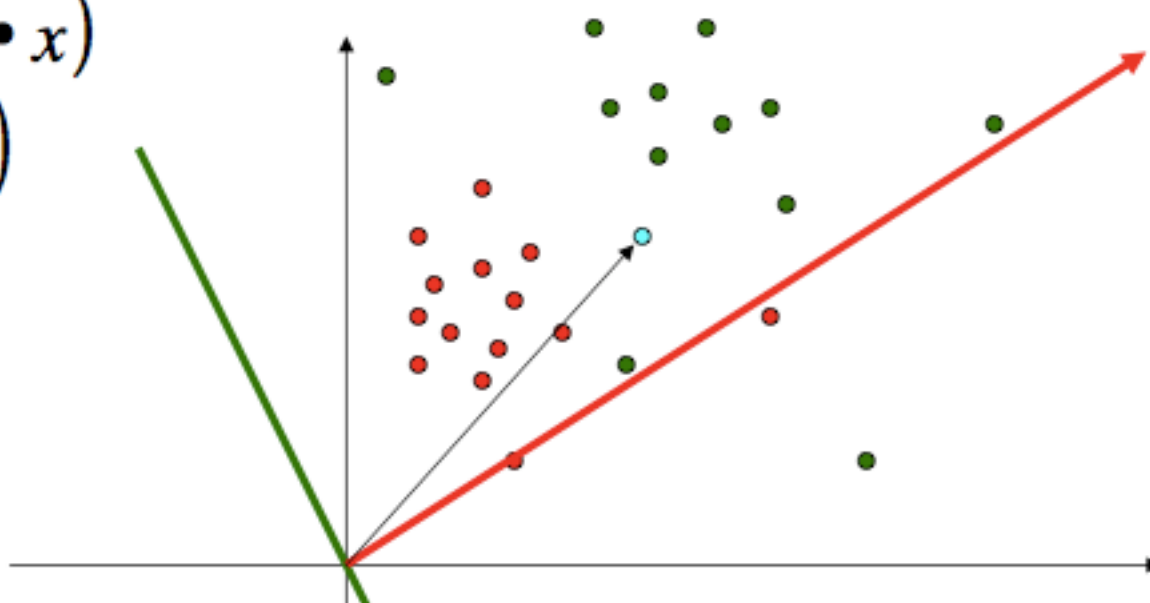
Perceptron Learning

- Operation is a dot product

$$y = \text{sgn}\left(\sum w_j x_j\right)$$

$$y = \text{sgn}(w \bullet x)$$

$$= \text{sgn}(w^T x)$$

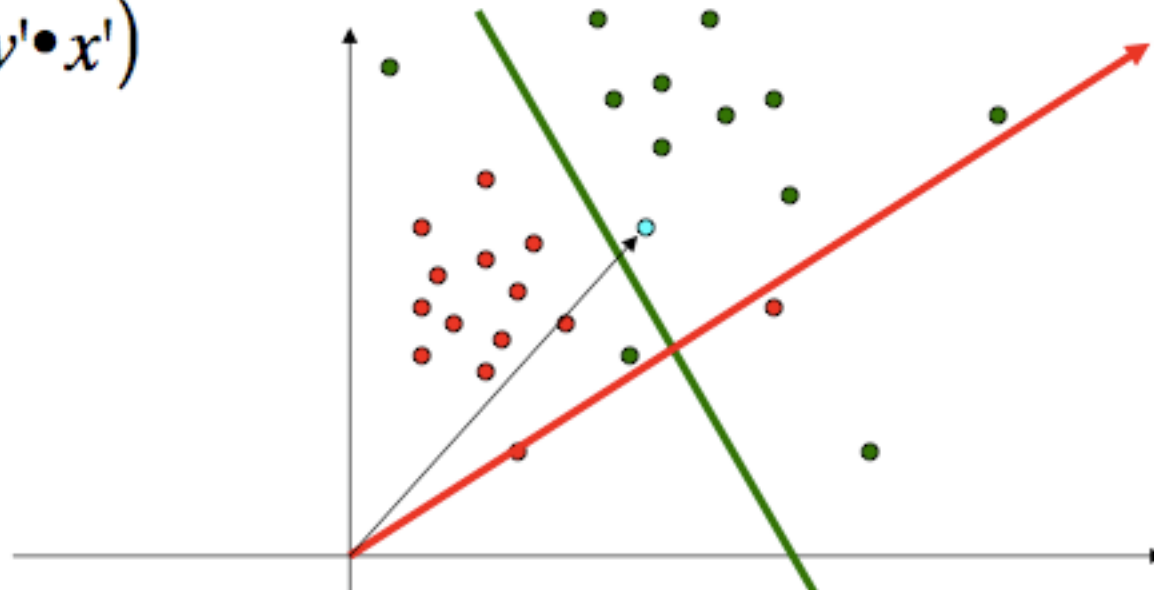


Perceptron Classification

- Adding bias term

$$y = \text{sgn}\left(\sum w_j x_j - b\right)$$

$$y = \text{sgn}(w' \bullet x')$$



Perceptron Learning

- Very simple rule:

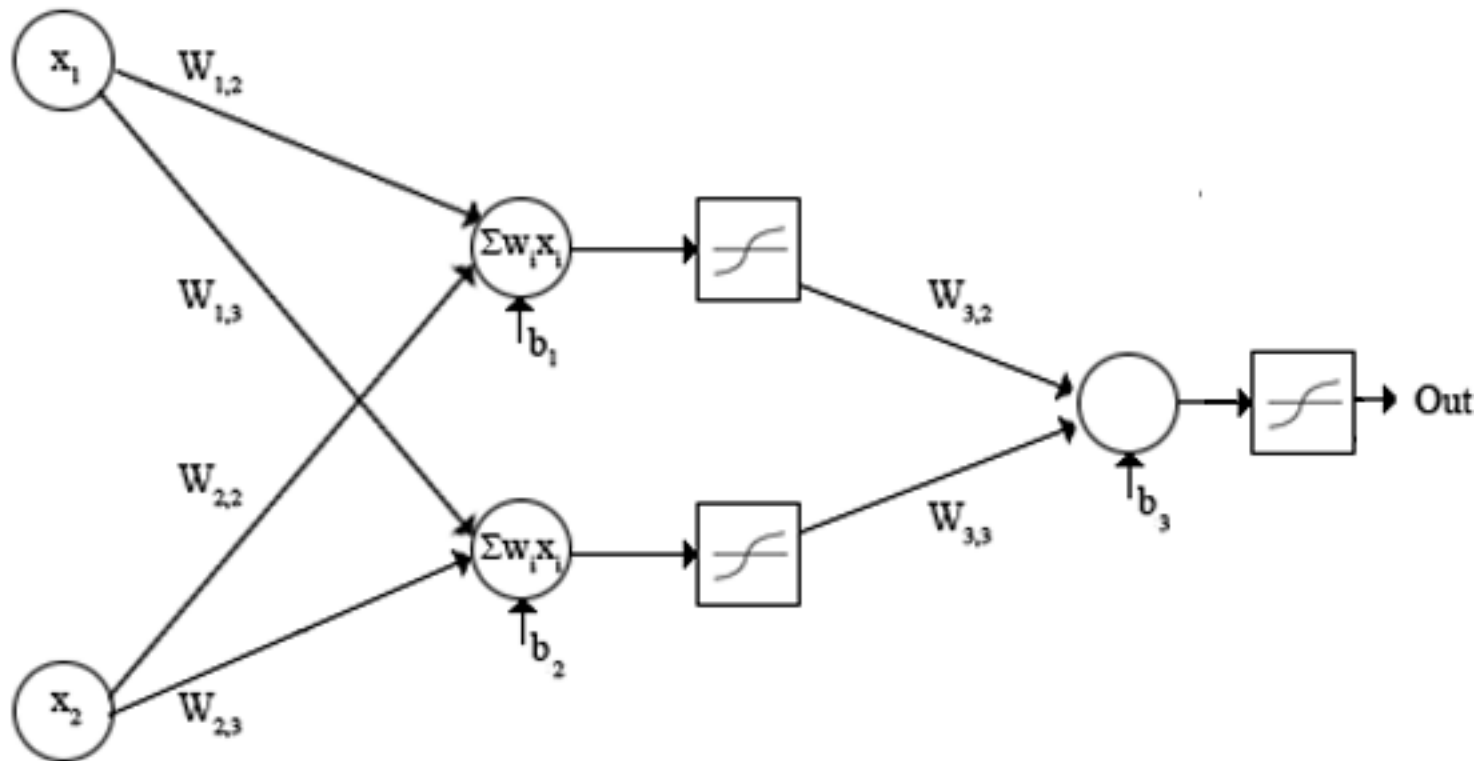
Difference in desired
and actual answer

$$w' = w + \alpha(y_i - \text{sgn}(w^T x_i))x_i$$

Alpha is a
learning rate

- This gives us a method of modifying the weights based on misclassifications.

Multilayer Perceptron



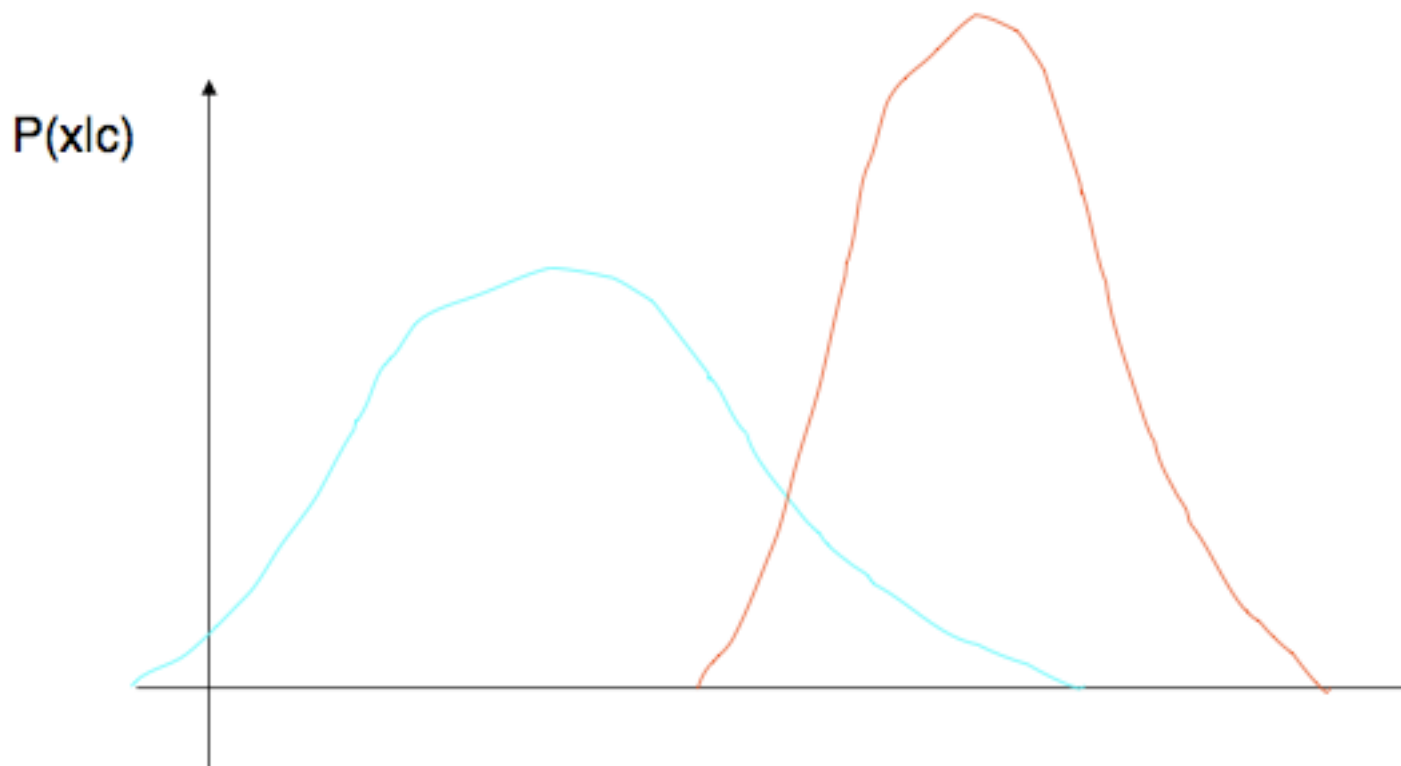
Naïve Bayes

- Another alternative is to look at the problem probabilistically
- Define class probabilities

$$P(x|c = 1), P(x|c = -1)$$

Consider a 1-D example with the distributions plotted....

Naïve Bayes



Best choice will be:

$$c^* = \arg \max_c P(c|x)$$

Naïve Bayes

Best choice will be:

$$c^* = \arg \max_c P(x|c)P(c)$$

Prior

- How does this work?
- Let's review some probability theory?

Axioms of Probability Theory

- All probabilities between 0 and 1

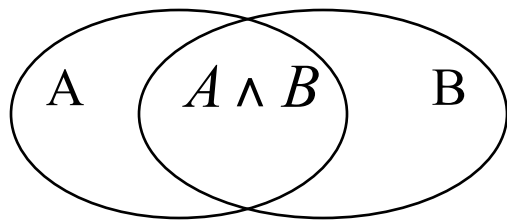
$$0 \leq P(A) \leq 1$$

- True proposition has probability 1, false has probability 0.

$$P(\text{true}) = 1 \quad P(\text{false}) = 0.$$

- The probability of disjunction is:

$$P(A \vee B) = P(A) + P(B) - P(A \wedge B)$$

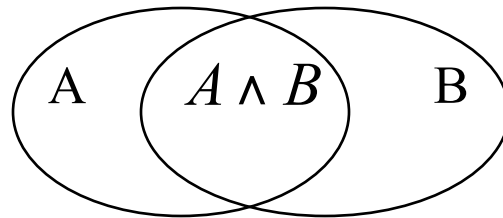


- A and B are independent if the joint probability $P(A \text{ and } B) = P(A)P(B)$

Conditional Probability

- $P(A | B)$ is the probability of A given B
- Assumes that B is all and only information known.
- Defined by:

$$P(A | B) = \frac{P(A \wedge B)}{P(B)}$$



Bayes Theorem

$$P(H | E) = \frac{P(E | H)P(H)}{P(E)}$$

Simple proof from definition of conditional probability:

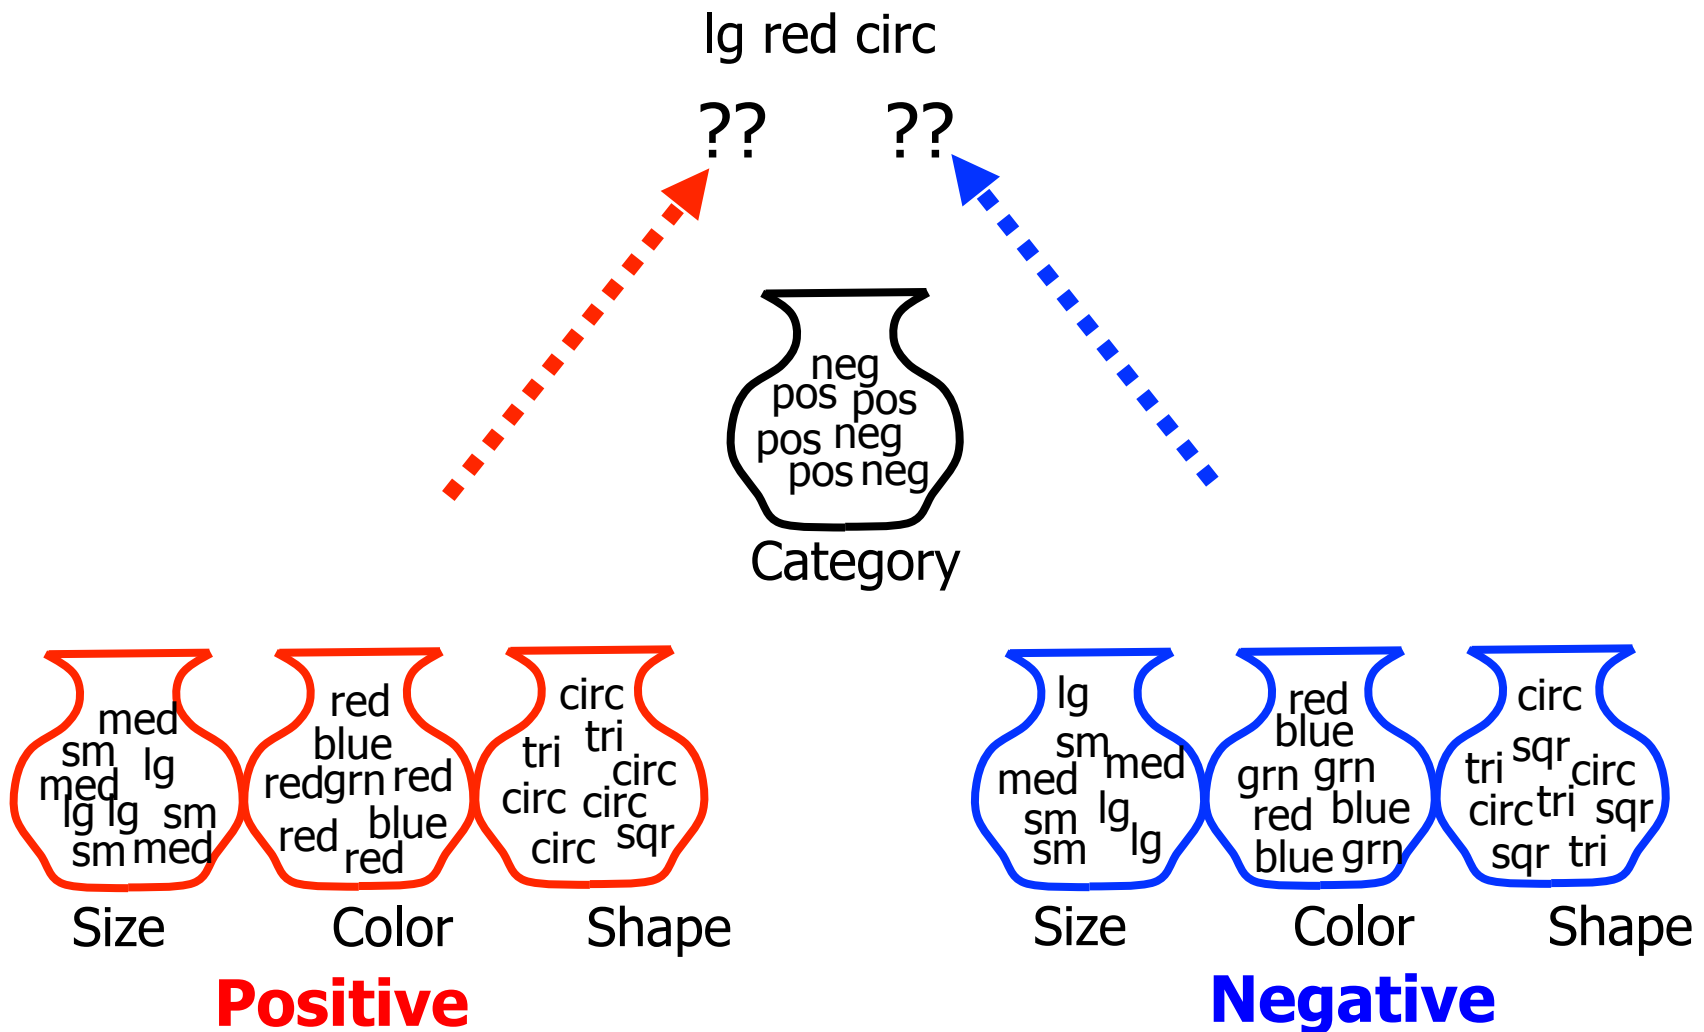
$$P(H | E) = \frac{P(H \wedge E)}{P(E)} \quad (\text{Def. cond. prob.})$$

$$P(E | H) = \frac{P(H \wedge E)}{P(H)} \quad (\text{Def. cond. prob.})$$

$$P(H \wedge E) = P(E | H)P(H)$$

$$\text{QED: } P(H | E) = \frac{P(E | H)P(H)}{P(E)}$$

Naïve Bayes Inference Problem



Naïve Bayesian Categorization

- If we assume features of an instance are independent **given the category** (***conditionally independent***).

$$P(X | Y) = P(X_1, X_2, \dots, X_n | Y) = \prod_{i=1}^n P(X_i | Y)$$

- Therefore, we then only need to know $P(X_i | Y)$ for each possible pair of a feature-value and a category.
- If Y and all X_i are binary, this requires specifying only $2n$ parameters:
 - $P(X_i=\text{true} | Y=\text{true})$ and $P(X_i=\text{true} | Y=\text{false})$ for each X_i
 - $P(X_i=\text{false} | Y) = 1 - P(X_i=\text{true} | Y)$
- Compared to specifying 2^n parameters without any independence assumptions.

Naïve Bayes Example

| Probability | positive | negative |
|-----------------------------|----------|----------|
| $P(Y)$ | 0.5 | 0.5 |
| $P(\text{small} \mid Y)$ | 0.4 | 0.4 |
| $P(\text{medium} \mid Y)$ | 0.1 | 0.2 |
| $P(\text{large} \mid Y)$ | 0.5 | 0.4 |
| $P(\text{red} \mid Y)$ | 0.9 | 0.3 |
| $P(\text{blue} \mid Y)$ | 0.05 | 0.3 |
| $P(\text{green} \mid Y)$ | 0.05 | 0.4 |
| $P(\text{square} \mid Y)$ | 0.05 | 0.4 |
| $P(\text{triangle} \mid Y)$ | 0.05 | 0.3 |
| $P(\text{circle} \mid Y)$ | 0.9 | 0.3 |

We learn these probabilities by employing frequency counting techniques on the training data.

Test Instance:
<medium ,red, circle>

Naïve Bayes Example

| Probability | positive | negative |
|------------------------|----------|----------|
| $P(Y)$ | 0.5 | 0.5 |
| $P(\text{medium} Y)$ | 0.1 | 0.2 |
| $P(\text{red} Y)$ | 0.9 | 0.3 |
| $P(\text{circle} Y)$ | 0.9 | 0.3 |

Test Instance:
<medium ,red, circle>

Answer:
Drawn from the positive urn

$$\begin{aligned} P(\text{positive} | X) &= P(\text{positive}) * P(\text{medium} | \text{positive}) * P(\text{red} | \text{positive}) * P(\text{circle} | \text{positive}) / P(X) \\ &= \frac{0.5 * 0.1 * 0.9 * 0.9}{0.0495} = 0.8181 \end{aligned}$$

$$\begin{aligned} P(\text{negative} | X) &= P(\text{negative}) * P(\text{medium} | \text{negative}) * P(\text{red} | \text{negative}) * P(\text{circle} | \text{negative}) / P(X) \\ &= \frac{0.5 * 0.2 * 0.3 * 0.3}{0.0495} = 0.1818 \end{aligned}$$

$$P(\text{positive} | X) + P(\text{negative} | X) = 0.8181 + 0.1818 = 1$$

$$P(X) = (0.0405 + 0.009) = 0.0495$$

For purposes of making a decision, we can ignore the denominator since it is the same for both Classes.

What can we do with learning?

- Learn parameters that we normally have to specify
- Find good color thresholds for our vision algorithms
- Use reinforcement learning to do maze solving or to identify good kicks for Robocup players
- Use genetic algorithms to find good controller parameters
- Use Bayesian reasoning for localization
- Also these algorithms can be combined in interesting ways.